

# Alternative to relational database system – NoSQL

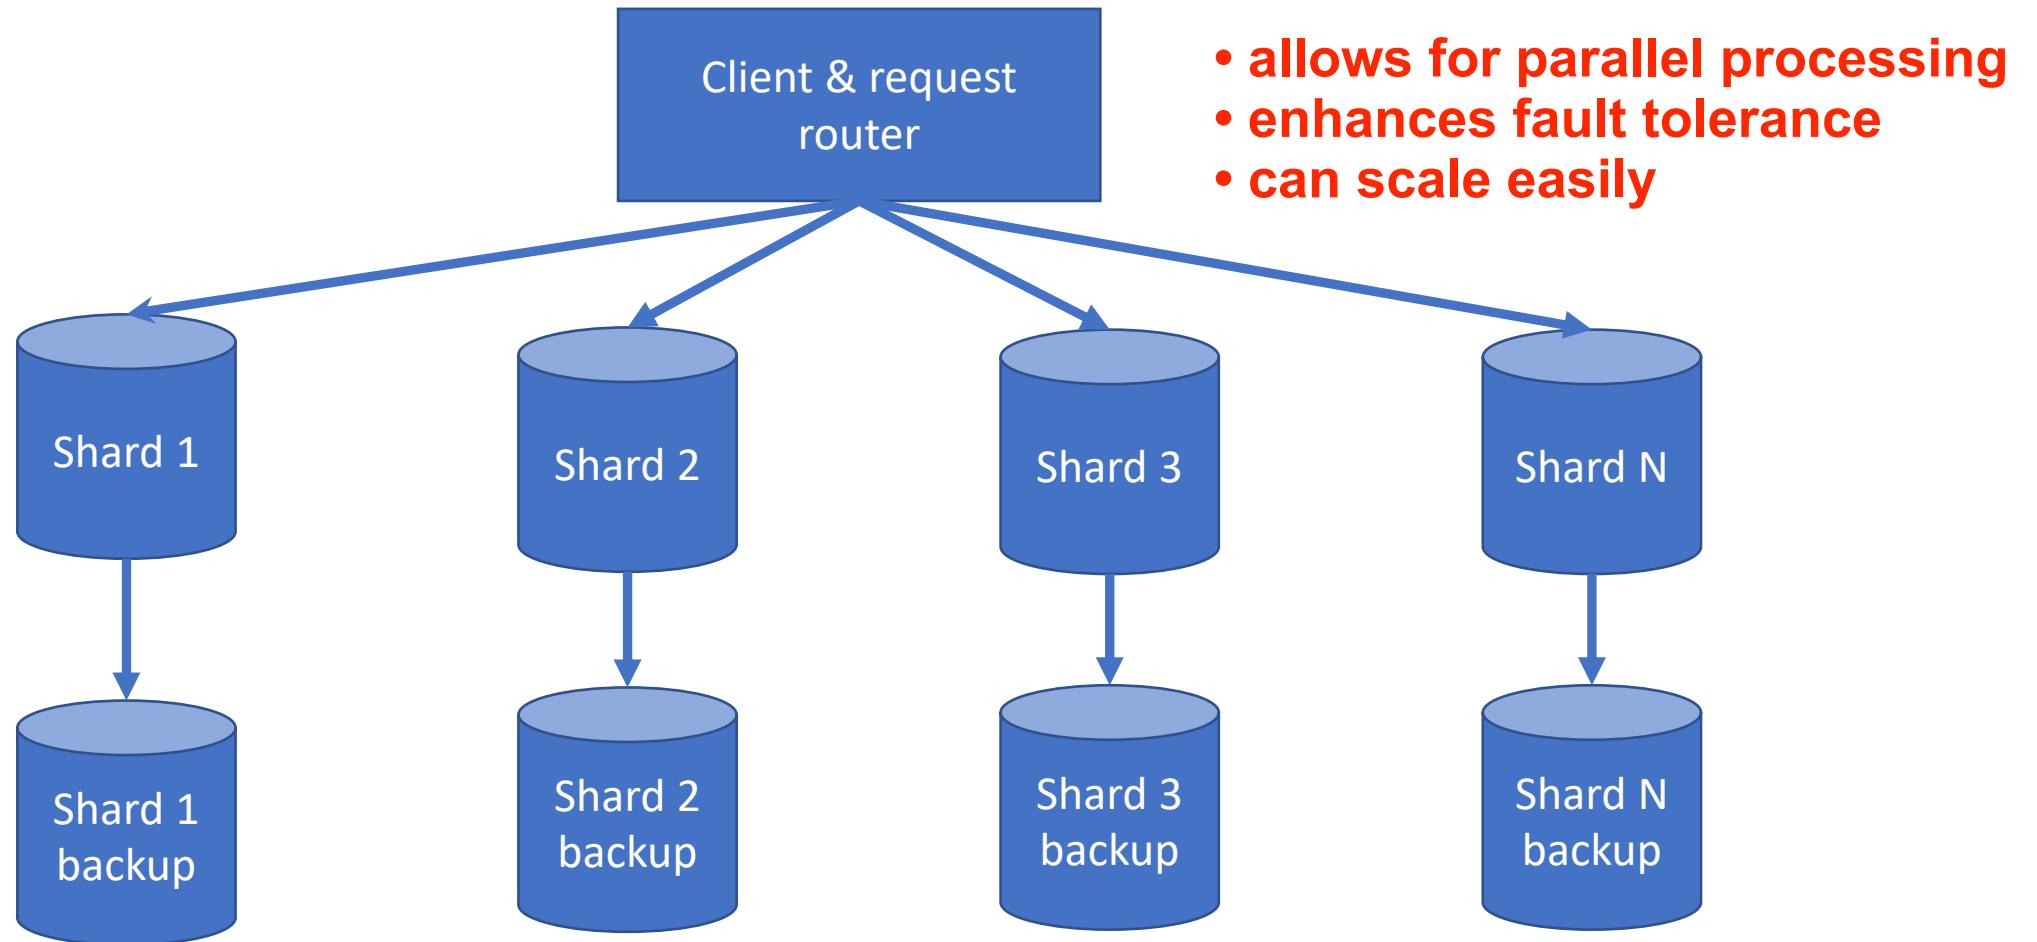
**Not Only SQL**

Bernard Lee Kok Bang

# Do we really need SQL?

- High-transaction queries are probably pretty simple once de-normalized [such as using JSON file format]
- A simple get/put API may meet our needs [*“given this customer ID, give me back this list of information”*]
- Looking up values for a given key is simple, fast, and scalable [at high level]
- Questions:
  - *What’s the item information associated with this item?*
  - *What pages have this customer looked at?*
- Don’t need a rich query language for answering the above question
- Don’t need big fancy relational database; instead, we need more scalable system that can easily horizontally partitioned for given key ranges

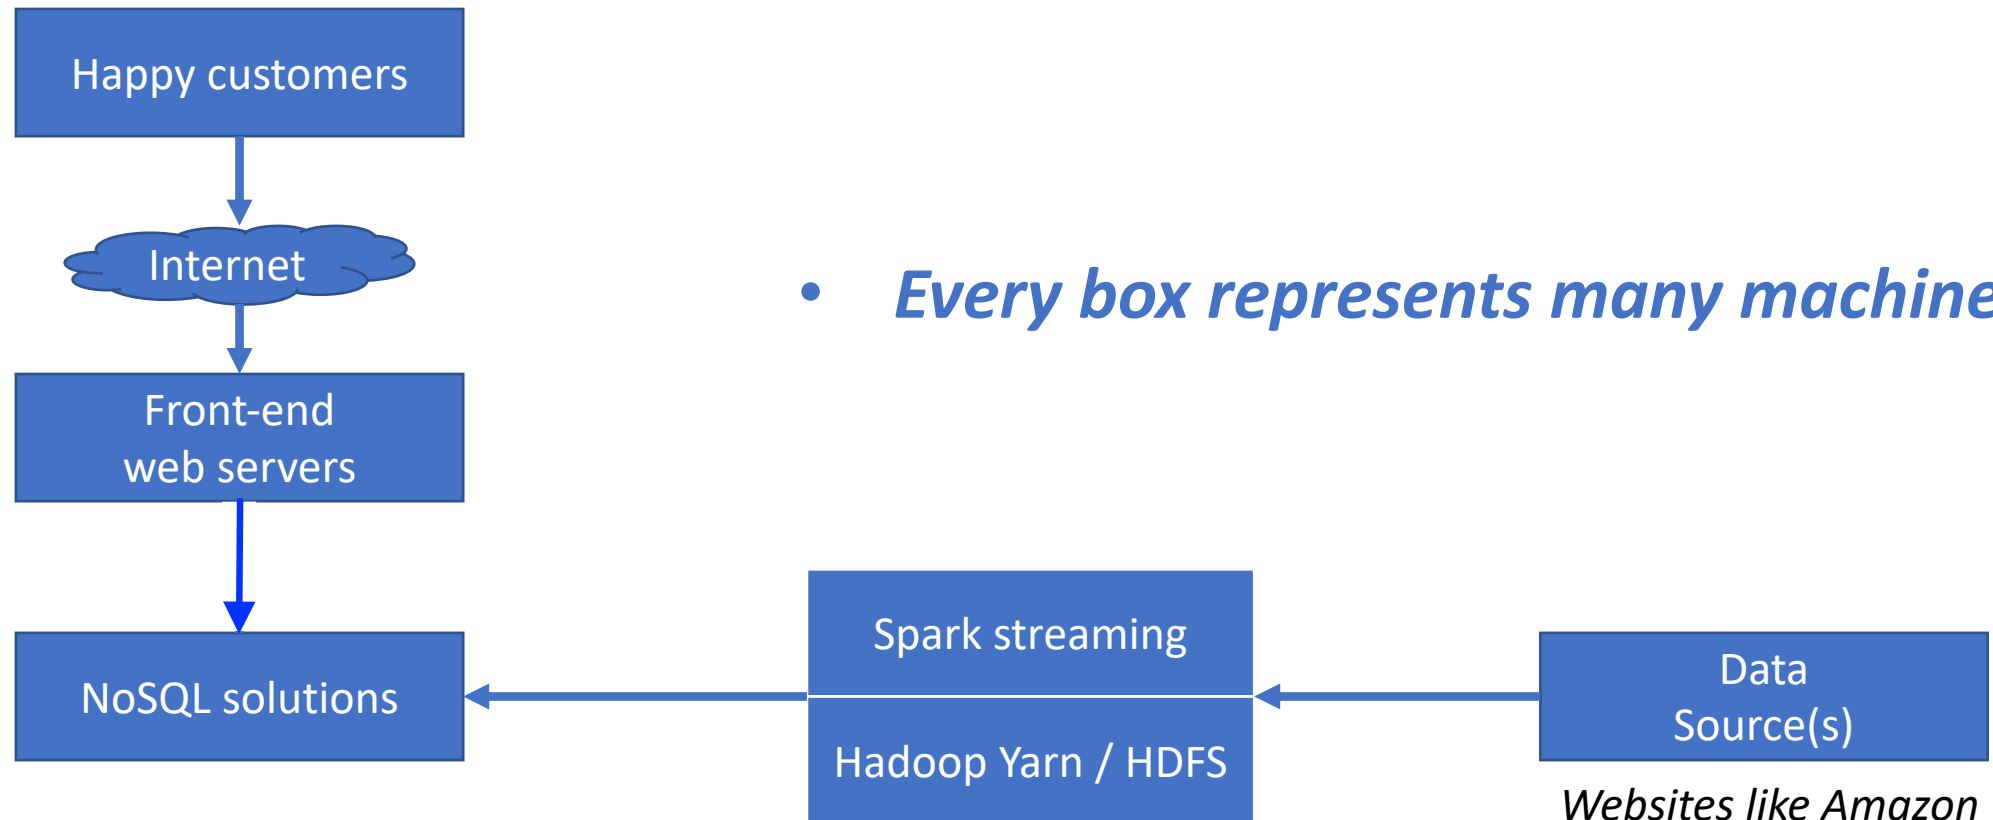
# NoSQL database sample architecture



# Use the right tool for the job

- For *analytic queries*, Hive, Pig, Spark, etc. work great [business report]
- Exporting data to MySQL is plenty fast for most applications too
- But if we work at *giant scale* [Amazon, Google...] - export our data to a non-relational database for fast and scalable serving of that data to web applications, etc.

# Sample application architecture (simplified)



# HBase

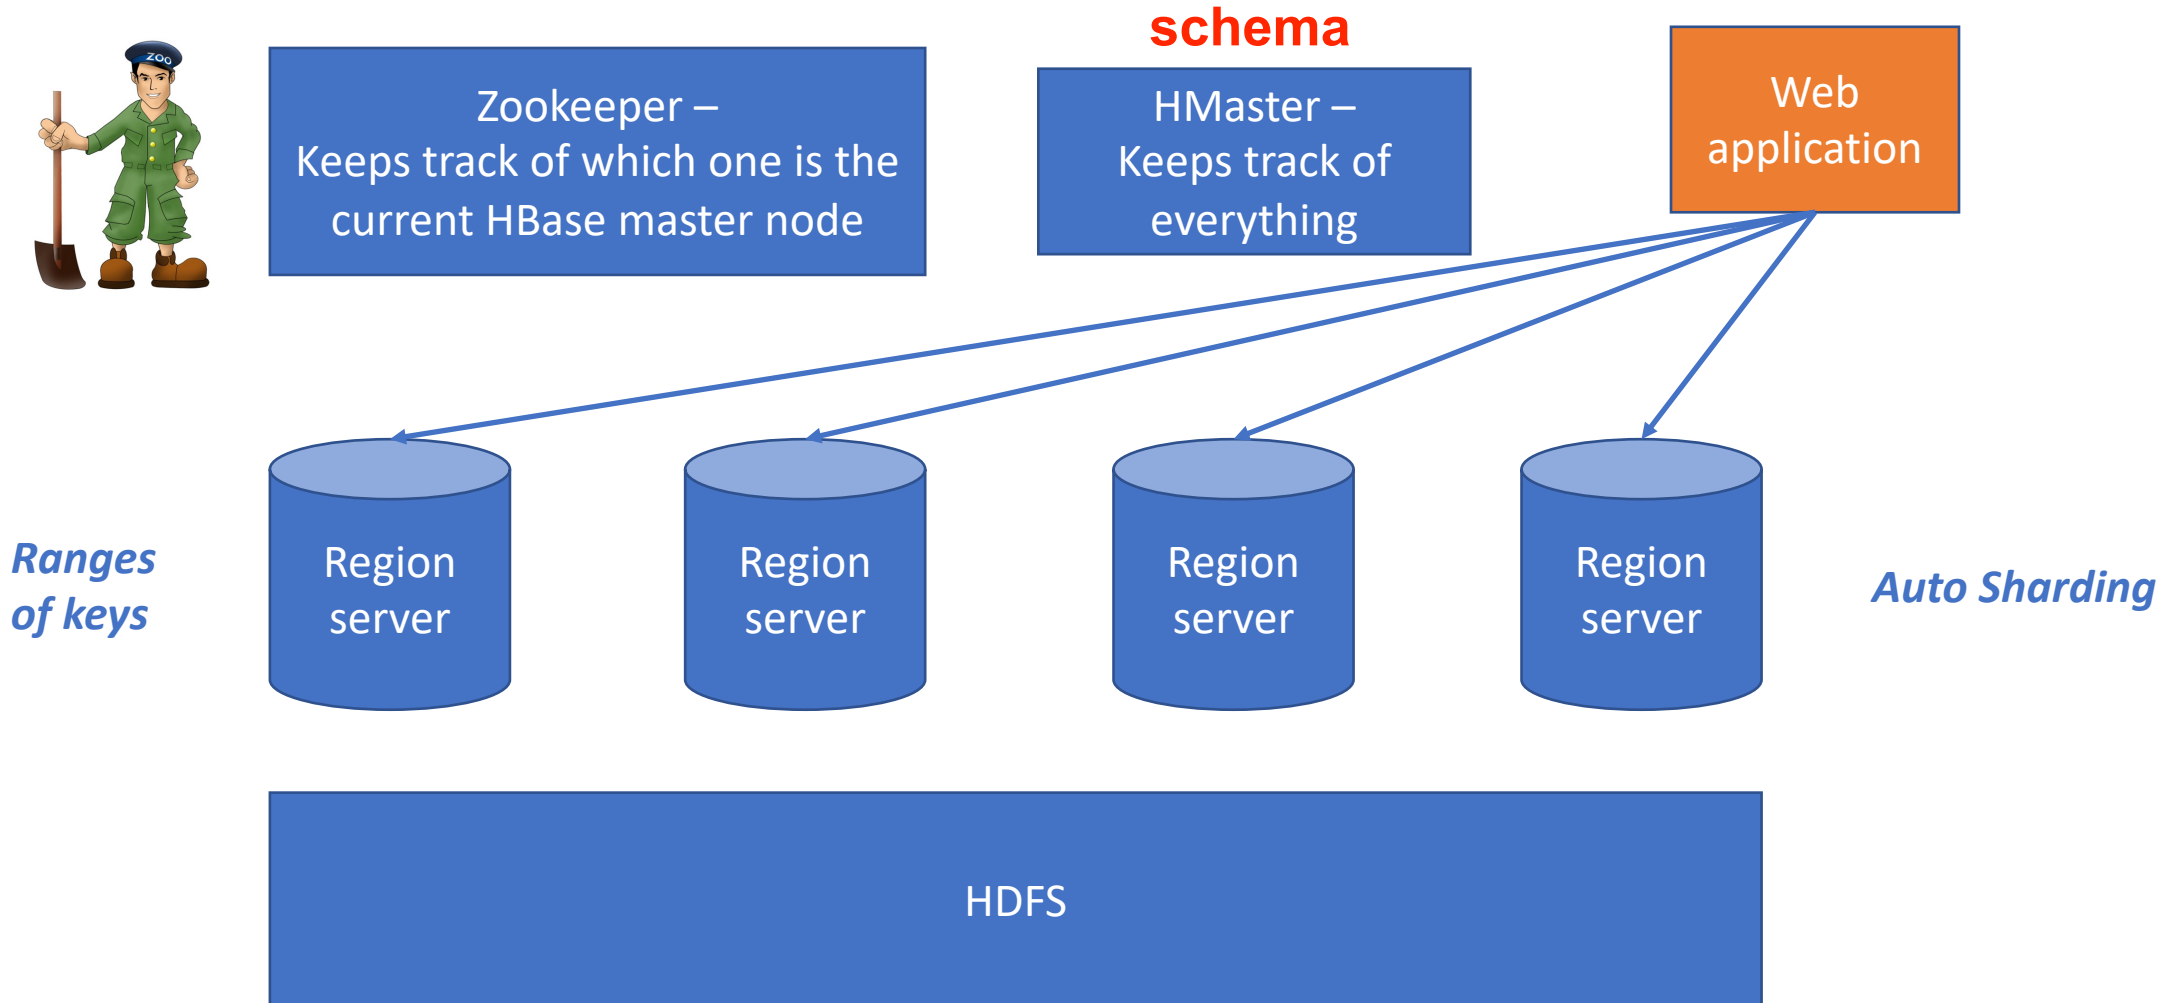
- Non-relational, scalable database built on HDFS
- Expose massive data sitting on HDFS to web service or web application [that can operate very quickly and high scale]
- Don't have a query language
- Have an API that can quickly answer the question
  - “what are the values for this key”
  - “store this value for this key”



# CRUD operations

- Create
- Read
- Update
- Delete
- No query language, only CRUD API
- Offers transactional access -> ***FAST & LARGE SCALE!!!***

# HBase architecture





# *HBase Data Model*

- HBase organizes data as a large, sparse, distributed table.
- Think of it like a massive spreadsheet where:
  - Rows can have millions of columns
  - Most cells are empty (sparse)
  - Columns don't need to be predefined
  - Data is stored across multiple machines

# *Key Components of HBase Data Model*

- *Row Key:*

- like a primary key in relational databases, but it's the only way to uniquely identify a row.
- Rows are automatically sorted by row key *lexicographically*.

- *Column Families:*

- groups of related columns that are stored together physically.
- must define column families when creating a table, but individual columns within families can be added dynamically.

- *Columns (Qualifiers):*

- Within each column family, we can have any number of columns.
- The full column name is "family:qualifier" (e.g., "*personal:name*").

# *"Lexicographical" in HBase*

## **Examples:**

If you have these row keys:

- "apple"
- "application"
- "app"
- "banana"
- "band"

They would be stored in this lexicographic order:

1. "app"
2. "apple"
3. "application"
4. "banana"
5. "band"

## **With numbers as strings:**

- "1" comes before "2"
- "10" comes before "2" (because it compares "1" vs "2" first)
- "100" comes before "11"

## **Practical implications:**

If you have employee IDs like "emp001", "emp002", "emp010", they'll be ordered:

- emp001
- emp002
- emp010

But if you used "emp1", "emp2", "emp10", they'd be ordered:

- emp1
- emp10
- emp2

# Using HBase Shell - PuTTY

*#Opens the HBase shell*

*hbase shell*

*#Lists down all the tables present in HBase*

*list*

*#Creates a new table*

*create 'newtbl', 'knowledge'*

*#Checks if the table was created*

*describe 'newtbl'*

*#Checks the status of HBase*

*status 'summary'*

*#Put some data into the 'newtbl'*

*put 'newtbl', 'r1', 'knowledge:sports', 'cricket'*

*put 'newtbl', 'r1', 'knowledge:science', 'chemistry'*

*put 'newtbl', 'r1', 'knowledge:science', 'physics'*

*put 'newtbl', 'r2', 'knowledge:economics', 'macro economics'*

*put 'newtbl', 'r2', 'knowledge:music', 'pop music'*

# Using HBase Shell – PuTTY (cont...)

*#List the contents of the 'newtbl'*  
*scan 'newtbl'*

*#Checks if the table is enabled*  
*is\_enabled 'newtbl'*

*#Disables the table*  
*disable 'newtbl'*

*#Lists the contents of the 'newtbl'. Note that this will throw an error as the table is disabled.*  
*scan 'newtbl'*

*#Updates column family in the table*  
*alter 'newtbl', 'test\_info'*

*#Enables the table*  
*enable 'newtbl'*

*#Checks the column families after updating*  
*describe 'newtbl'*

*#Extracts the values for r1 in the 'newtbl'*  
*get 'newtbl', 'r1'*

# Using HBase Shell – PuTTY (cont...)

*#Adds new information to r1 for economics. Note that this will update the table but will not override the information*

*put 'newtbl', 'r1', 'knowledge:economics', 'market economics'*

*#Displays the results for r1*

*get 'newtbl', 'r1'*

*#Count number of rows in 'newtbl'*

*count 'newtbl'*

*# Delete the 'newtbl'. Note: Error because the 'newtbl' is still enabled*

*drop 'newtbl'*

*#Disable first, then only delete the the 'newtbl'*

*disable 'newtbl'*

*# Check if the 'newtbl' still available*

*list*

What if I  
wanted to add  
in more values  
in same row?

```
▶ # Disable the table if it exists
disable 'newtbl'

# Alter the table to keep up to 3 versions of each cell
alter 'newtbl', NAME => 'knowledge', VERSIONS => 3

# Enable the table
enable 'newtbl'

# Put the values
put 'newtbl', 'r1', 'knowledge:sports', 'cricket'
put 'newtbl', 'r1', 'knowledge:sports', 'football'
put 'newtbl', 'r1', 'knowledge:sports', 'basketball'

# Retrieve the versions
get 'newtbl', 'r1', {COLUMN => 'knowledge:sports', VERSIONS => 3}
```

# Using HBase Shell – PuTTY (cont...)

*#Check current HBase user*  
*whoami*

*#Different HBase table command usages and its syntaxes*  
*table\_help*


*#HBase status*  
*status*

*#HBase version*  
*version*




# Let's play with HBase

- Create a HBase table for *movie ratings grouped by user*
- Show how we can quickly query for individual users
- Good example of sparse data




# How are we doing it?



# REST API

## What is Rest API?



- providing standards between computer systems on the web, making it easier for systems to communicate with each other

<https://bit.ly/3whbwMF>

# Change some settings: Port 8000 & Ambari

- Open a port so that Python script can talk to REST service
  1. Right click Horthonworks Docker Sandbox on VirtualBox
  2. Choose Settings, Network, Advanced, Port Forwarding
  3. Click the ADD button and add in the following details  
→ Name: *HBase REST*; Protocol: *TCP*; Host IP: *127.0.0.1*; Host Port: *8000*; Guest Port: *8000*
  4. If Port 8000 has been occupied, then just change the Name to *HBase REST*
  5. Log in to Ambari using *admin* username
  6. Choose *HBase* [left panel], and *start* the HBase under *Service Action*

# Change some settings: PuTTY

**Starts** REST server sitting on top of HBase and HDFS

1. Login as maria\_dev

*su root*

*Password:*

*/usr/hdp/current/hbase-master/bin/hbase-daemon.sh start rest -p 8000 --infoport 8001*

# Start coding

```
pip install starbase
```

```
from starbase import Connection
```

```
c = Connection("127.0.0.1", "8000")
```

```
ratings = c.table('ratings')
```

```
if (ratings.exists()):
```

```
    print ("Dropping existing ratings table\n")
```

```
    ratings.drop()
```

```
ratings.create('rating')
```

```
print("Parsing the ml-100k ratings data... \n")
```

```
ratingFile = open(r"c:/Downloads/ml-100k/ml-100k/u.data", "r")
```

```
batch = ratings.batch()
```

```
for line in ratingFile:
```

```
    (userID, movieID, rating, timestamp) = line.split()  
    batch.update(userID, {'rating': {movieID: rating}})
```

```
ratingFile.close()
```

```
print("Committing ratings data to HBase via REST service\n")
```

```
batch.commit(finalize=TRUE)  
True
```

```
print("Get back ratings for some users...\n")
```

```
print("Ratings for user ID 33:\n")
```

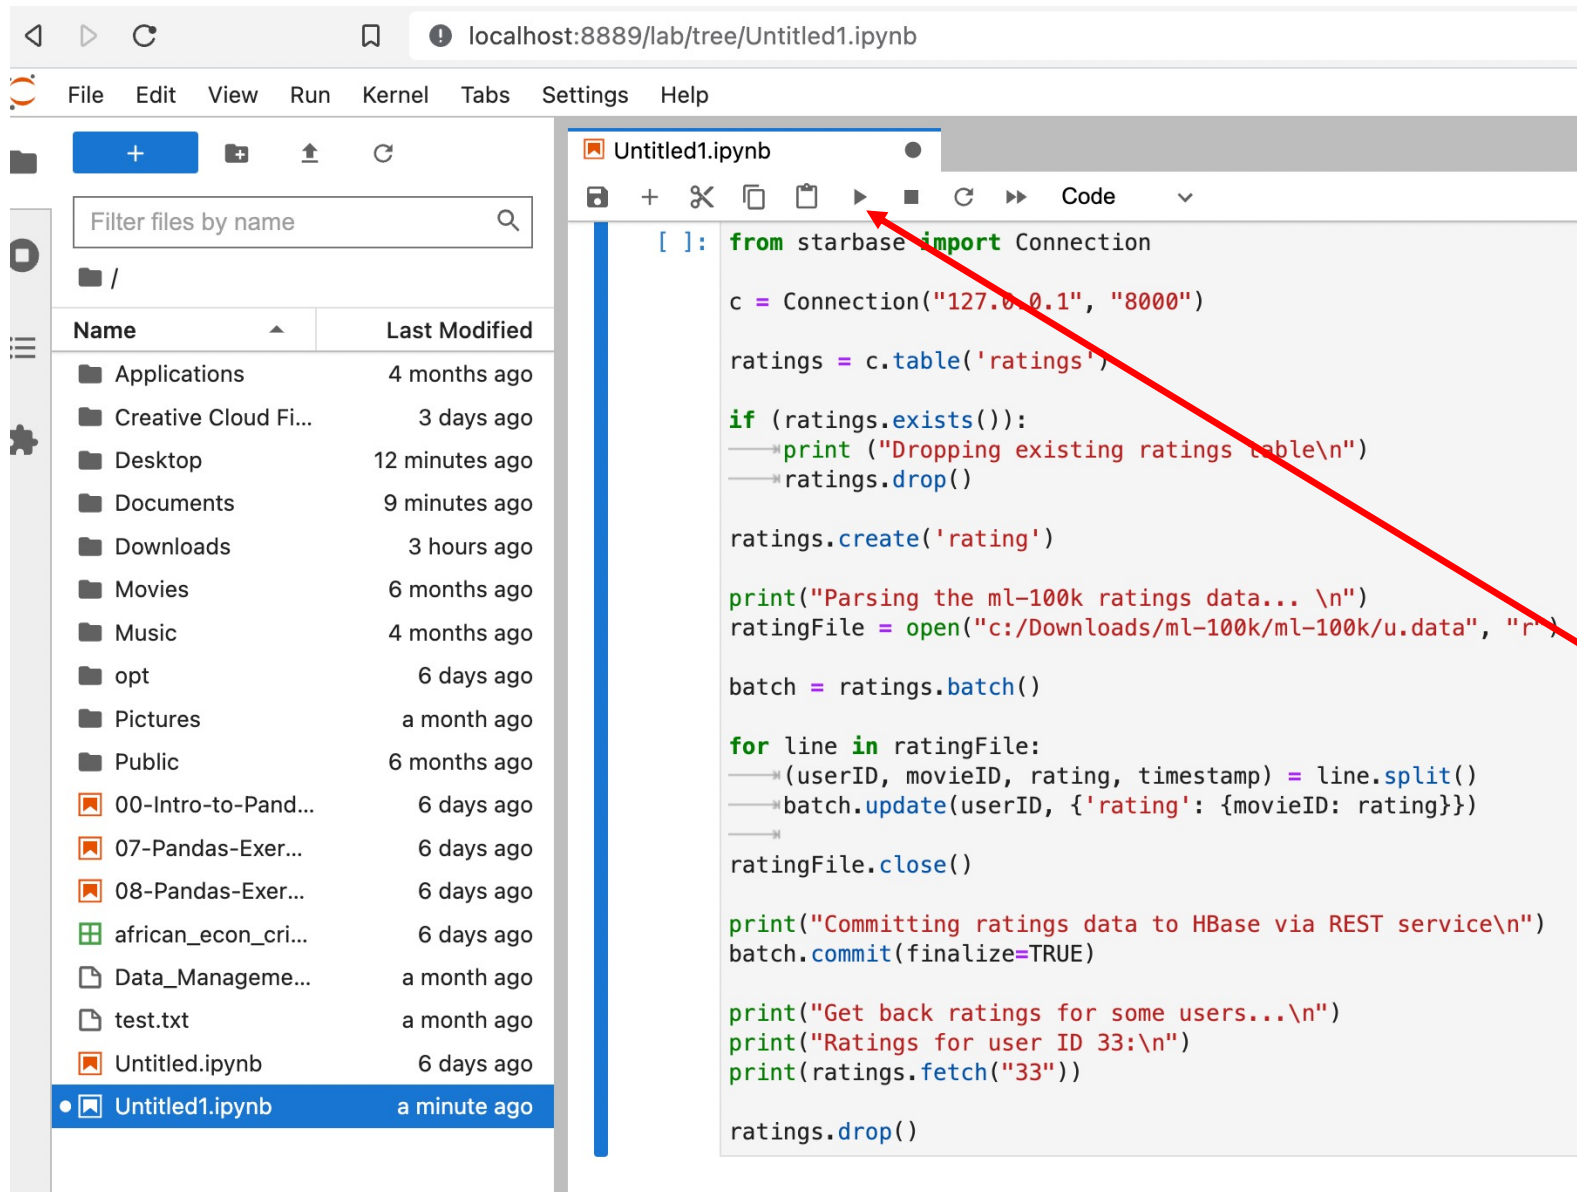
```
print(ratings.fetch("33"))
```

```
ratings.drop()
```



***Change directory accordingly***

# Using JupyterLab



The screenshot shows the JupyterLab web interface in a browser at `localhost:8889/lab/tree/Untitled1.ipynb`. The interface has a top menu bar with 'File', 'Edit', 'View', 'Run', 'Kernel', 'Tabs', 'Settings', and 'Help'. On the left is a file browser with a search bar and a table of files. The file `Untitled1.ipynb` is selected. The main area is a code editor for `Untitled1.ipynb`, showing Python code. A red arrow points to the 'Run' button (a play icon) in the code editor's toolbar.

File browser contents:

Name	Last Modified
Applications	4 months ago
Creative Cloud Fi...	3 days ago
Desktop	12 minutes ago
Documents	9 minutes ago
Downloads	3 hours ago
Movies	6 months ago
Music	4 months ago
opt	6 days ago
Pictures	a month ago
Public	6 months ago
00-Intro-to-Pand...	6 days ago
07-Pandas-Exer...	6 days ago
08-Pandas-Exer...	6 days ago
african_econ_cri...	6 days ago
Data_Manageme...	a month ago
test.txt	a month ago
Untitled.ipynb	6 days ago
Untitled1.ipynb	a minute ago

```
[ ]: from starbase import Connection

c = Connection("127.0.0.1", "8000")

ratings = c.table('ratings')

if (ratings.exists()):
    print("Dropping existing ratings table\n")
    ratings.drop()

ratings.create('rating')

print("Parsing the ml-100k ratings data... \n")
ratingFile = open("c:/Downloads/ml-100k/ml-100k/u.data", "r")

batch = ratings.batch()

for line in ratingFile:
    (userID, movieID, rating, timestamp) = line.split()
    batch.update(userID, {'rating': {movieID: rating}})

ratingFile.close()

print("Committing ratings data to HBase via REST service\n")
batch.commit(finalize=TRUE)

print("Get back ratings for some users...\n")
print("Ratings for user ID 33:\n")
print(ratings.fetch("33"))

ratings.drop()
```

**Click the icon  
to start running the  
python script**

# Output

```
Python
'185': '4', '188': '3', '189': '3', '4': '3', '97': '3', '6': '5', '94': '2', '99': '3',
'228': '5', '227': '4', '165': '5', '166': '5', '224': '5', '223': '5', '222': '4', '221':
'5', '12': '5', '15': '5', '14': '5', '17': '3', '16': '5', '19': '5', '18': '4', '272':
'5', '153': '3', '152': '5', '155': '2', '154': '5', '157': '4', '156': '4', '159': '3',
'239': '4', '83': '3', '234': '4', '235': '5', '236': '4', '237': '2', '230': '4', '231':
'3', '46': '4', '47': '4', '44': '5', '45': '5', '42': '5', '43': '4', '40': '3', '118':
'146': '4', '200': '3', '203': '4', '202': '5', '205': '3', '204': '5', '207': '5', '206':
'149': '2', '77': '4', '76': '4', '75': '4', '74': '1', '73': '3', '72': '4', '71': '3',
'2': '3', '263': '1', '262': '3', '261': '1', '41': '2', '260': '1', '267': '4', '67': '3
Ratings for user ID 33:
{'rating': {'751': '4', '880': '3', '339': '3', '895': '3', '313': '5', '872': '3', '333':
'258': '4', '260': '4', '678': '4', '328': '4', '288': '4', '292': '4', '879': '3', '300':
'329': '4', '348': '4', '682': '4'}}
In [2]:
```

*Ratings for  
user ID 33*

**To stop REST interface:  
/usr/hdp/current/hbase-master/bin/hbase-daemon.sh stop rest**